

FinalProject_SDA.ipynb

September 7, 2018

0.0.1 Jake Dineen and Mason David

0.0.2 Scripting For Data Analysis

0.0.3 Final Project

Preface This project will center around collections of baseball data ranging from 1871 - 2014, essentially encompassing all recorded statistics over the course of the history of Major League Baseball. Sean Lahman and a group of researchers are responsible for the collection and storage of most of the main files we will be working with, although there have been some crowdsourcing attempts at expanding the original sets to include such things as college statistics.

Data All data is available to us via a github repository. To explore some of text mining/web scraping tactics, and to exemplify some additional learned skills, we will aggregate the CSV links in using urllib and BeautifulSoup. This will be an exercise of munging semi-structured HTML while the rest of our task involves structured, tabular data. As seen below, there are 27 different data sets available to us, but not all of them will be relevant for our project's goals.

Our Goal(s) We have a number of goals associated with the project: - Read in data directly from the web - Parse CSV links from HTML and automate a process for storing data. - Work to map datasets together based on certain metadata. The CSV files work similar to a sql db, so we'll need to merge on keys/indices. - Summarize descriptive statistics in a variety of ways: - Print out single season record holders for major statistics, along with the corresponding metrics - Print out All Time record holders for major statistics, along with the corresponding metrics - Visualize and describe through tables/charts the way statistics and salaries have changed over time * This work will likely focus on the use of pandas and numpy - Potentially roll this up into a mini chatbot that requests a user input and outputs a result. - Create a hierarchical binary classification problem centered on hall of fame candidacy: - Data munging for cleanup, as well as feature distributions. - Statistical Measures (Chi Square/Correlation/Multicollinearity) for dimensionality reduction - Function Approximation via Gradient Boosted Trees/Logit/Random Forests - Utilize a sigmoid or softmax activation function to output probability distribution for class mapping. - Using an interpretable model to understand feature importance/ranking/information gain. - Display what drives hall of fame potential with a meaningful metric associated with it. - Potentially roll this up into a mini chatbot that requests a user input (statistics) and outputs the probability that they will make it to the Hall of Fame.

Store data sources in Variables

```
In [6]: #link to gdoc: https://docs.google.com/document/d/1PJ1KdNt6EOcvMgaZidQbFdKrhna1x2JuQxE
#https://stackoverflow.com/questions/15517483/how-to-extract-urls-from-an-html-page-in
url = 'https://github.com/chadwickbureau/baseballdatabank/tree/master/core'
base_url_csv = 'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/mast
```

Extract information from HTML page

- The goal here is to pull out all of the CSVs that we'll eventually read into a structured format. It could have been done manually, but shows an understanding of webcrawling and HTML scraping. The final variable has a list of stored links containing all of the data in csv format.

```
In [7]: #from utils import get_csv_links
import urllib
from bs4 import BeautifulSoup
from IPython.core.display import display
import warnings
warnings.filterwarnings("ignore")

#from utils import get_csv_links. Stored in current working dir
# ^^ Wrap this up in a separate .py file and import it. Helper functions/Wrappers to b

def get_csv_links(url, base_url_csv):
    print('Scraping Webpage:', url)
    link_storage = []
    html = urllib.request.urlopen(url)
    html = html.read()
    soup = BeautifulSoup(html, "lxml")
    links = soup.find_all('a') #Find all href
    #Introduce logic to parse links containing .csv
    try:
        for tag in links:
            link = tag.get('href', None)
            if link is not None:
                if '.csv' in link:
                    link_storage.append(link)
    clean_urls = [] #instantiate
    #Looping through and adding the root URL
    clean_links = []
    #Need to extract after the last backslash
    for i in link_storage:
        clean_links.append(i.rsplit('/', 1)[-1])
    #Piecing URL strings together
    for i in clean_links:
        clean_urls.append(base_url_csv + i)
    print('Number of datasets: {}'.format(len(clean_urls)))
    return clean_urls
```

```

except:
    print('Something went wrong')

```

```

#URL is the page that we're scraping. Base URL is the root of the final URL
data_links = get_csv_links(url, base_url_csv)
data_links

```

Scraping Webpage: <https://github.com/chadwickbureau/baseballdatabank/tree/master/core>
Number of datasets: 27

```

Out[7]: ['https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Allstar',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Appear',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/AwardsS',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/AwardsS',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/AwardsS',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/AwardsS',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Batting',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Batting',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Colleg',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Fieldin',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Fieldin',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Fieldin',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/HallOfF',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/HomeGar',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Manag',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Manag',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Parks.',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/People',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Pitchin',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Pitchin',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Salari',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/School',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Series',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Teams.',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/TeamsF',
'https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/TeamsH

```

```

In [83]: #Need positions to act as a heirarchical filter
import pandas as pd
position = pd.read_csv('https://raw.githubusercontent.com/chadwickbureau/baseballdatabank/master/core/Positions.csv')
position.columns
positions= position[['playerID', 'POS']]
positions.drop_duplicates(subset = 'playerID', inplace= True)

```

Reading in all necessary data & Munging

```
In [84]: #Reading in batting csv
import pandas as pd
batting = pd.read_csv(data_links[6]) #index 6
batting.drop_duplicates(inplace = True) #No duplicate playerIDs
print('Batting dataset is {} rows by {} columns'.format(
    batting.shape[0], batting.shape[1]))
batting.head(5)
```

Batting dataset is 104324 rows by 22 columns

```
Out[84]:
```

	playerID	yearID	stint	teamID	lgID	G	AB	R	H	2B	...	RBI	SB	\
0	abercda01	1871	1	TR0	NaN	1	4	0	0	0	...	0.0	0.0	
1	addybo01	1871	1	RC1	NaN	25	118	30	32	6	...	13.0	8.0	
2	allisar01	1871	1	CL1	NaN	29	137	28	40	4	...	19.0	3.0	
3	allisdo01	1871	1	WS3	NaN	27	133	28	44	10	...	27.0	1.0	
4	ansonca01	1871	1	RC1	NaN	25	120	29	39	11	...	16.0	6.0	

	CS	BB	SO	IBB	HBP	SH	SF	GIDP
0	0.0	0	0.0	NaN	NaN	NaN	NaN	0.0
1	1.0	4	0.0	NaN	NaN	NaN	NaN	0.0
2	1.0	2	5.0	NaN	NaN	NaN	NaN	1.0
3	1.0	0	2.0	NaN	NaN	NaN	NaN	0.0
4	2.0	2	1.0	NaN	NaN	NaN	NaN	0.0

[5 rows x 22 columns]

```
In [10]: #Reading in pitching csv
pitching = pd.read_csv(data_links[-8]) #index 8
pitching.drop_duplicates(inplace = True) #no dupes
print('Pitching dataset is {} rows by {} columns'.format(
    pitching.shape[0], pitching.shape[1]))
pitching.head(5)
```

Pitching dataset is 45806 rows by 30 columns

```
Out[10]:
```

	playerID	yearID	stint	teamID	lgID	W	L	G	GS	CG	...	IBB	WP	\
0	bechtge01	1871	1	PH1	NaN	1	2	3	3	2	...	NaN	7	
1	brainas01	1871	1	WS3	NaN	12	15	30	30	30	...	NaN	7	
2	fergubo01	1871	1	NY2	NaN	0	0	1	0	0	...	NaN	2	
3	fishch01	1871	1	RC1	NaN	4	16	24	24	22	...	NaN	20	
4	fleetfr01	1871	1	NY2	NaN	0	1	1	1	1	...	NaN	0	

	HBP	BK	BFP	GF	R	SH	SF	GIDP
0	NaN	0	146.0	0	42	NaN	NaN	NaN
1	NaN	0	1291.0	0	292	NaN	NaN	NaN
2	NaN	0	14.0	0	9	NaN	NaN	NaN
3	NaN	0	1080.0	1	257	NaN	NaN	NaN

```
4 NaN 0 57.0 0 21 NaN NaN NaN
```

```
[5 rows x 30 columns]
```

```
In [11]: #Need people csv for mapping to player names
people = pd.read_csv(data_links[-9])
print('People dataset is {} rows by {} columns'.format(people.shape[0],
                                                         people.shape[1]))

people.head(5)
```

```
#This is where we pull in/map player names form.
```

People dataset is 19598 rows by 24 columns

```
Out[11]:
```

	playerID	birthYear	birthMonth	birthDay	birthCountry	birthState	\
0	aardsda01	1981.0	12.0	27.0	USA	CO	
1	aaronha01	1934.0	2.0	5.0	USA	AL	
2	aaronto01	1939.0	8.0	5.0	USA	AL	
3	aasedo01	1954.0	9.0	8.0	USA	CA	
4	abadan01	1972.0	8.0	25.0	USA	FL	

	birthCity	deathYear	deathMonth	deathDay	...	nameLast	\
0	Denver	NaN	NaN	NaN	...	Aardsma	
1	Mobile	NaN	NaN	NaN	...	Aaron	
2	Mobile	1984.0	8.0	16.0	...	Aaron	
3	Orange	NaN	NaN	NaN	...	Aase	
4	Palm Beach	NaN	NaN	NaN	...	Abad	

	nameGiven	weight	height	bats	throws	debut	finalGame	retroID	\
0	David Allan	215.0	75.0	R	R	2004-04-06	2015-08-23	aardd001	
1	Henry Louis	180.0	72.0	R	R	1954-04-13	1976-10-03	aaroh101	
2	Tommie Lee	190.0	75.0	R	R	1962-04-10	1971-09-26	aarot101	
3	Donald William	190.0	75.0	R	R	1977-07-26	1990-10-03	aased001	
4	Fausto Andres	184.0	73.0	L	L	2001-09-10	2006-04-13	abada001	

	bbrefID
0	aardsda01
1	aaronha01
2	aaronto01
3	aasedo01
4	abadan01

```
[5 rows x 24 columns]
```

```
In [85]: #Merging batting and people on playerID. This gives player metadata
import os
import sys
current_working_dir = os.getcwd() #Get working dir
```

```

batting_merged = pd.merge(
    batting, people, how='left', on='playerID') #pandas merge

batting_merged = pd.merge(batting_merged, positions, on = 'playerID')

batting_cols = batting_merged.columns #store var for column names. Will need for dim

#Save csv in working dir
batting_merged_filename = 'batting_merged.csv'
batting_merged.to_csv(batting_merged_filename)
#Print Logs
print('File saved to {}'.format(current_working_dir))
print('File name: {}'.format(batting_merged_filename))
print('{} has {} rows and {} columns'.format(
    batting_merged_filename, batting_merged.shape[0], batting_merged.shape[1]))
print('-----')
print('\033[1m' + 'NA Count Distribution by Column' +
    '\033[0m') #Display text output
print(batting_merged.isna().sum(axis=0))

#We want to get rid of all of the observations where a given name is not present.
print('-----')
print(
    '\033[1m Dropped {} observations where nameGiven was not registered \033[1m'.
    format(batting_merged['nameGiven'].isna().sum(axis=0)))

#Create a column that merges nameFirst and nameLast
batting_merged[
    'FullName'] = batting_merged['nameFirst'] + ' ' + batting_merged['nameLast']

#Want to drop some unnecessary columns
cols_to_drop = [
    'birthMonth', 'birthDay', 'birthState', 'birthCity', 'deathYear',
    'deathMonth', 'deathDay', 'deathState', 'deathCountry', 'deathCity',
    'finalGame', 'retroID', 'bbrefID'
]
print('-----')
print('Dropping Unnecessary Columns')
batting_merged.drop(columns=cols_to_drop, inplace=True)
print('-----')
#print('Filtering on Position Players Only')
#batting_merged = batting_merged[batting_merged['POS'] != 'P']
print('-----')
print('Displaying Cleaned Batting Data')
display(batting_merged.head(5)) #Display tail of df

```

File saved to C:\Users\jdine\Desktop\Syracuse\Term 7\Scripting for Data Analysis\FinalProject:
File name: batting_merged.csv

batting_merged.csv has 104113 rows and 46 columns

NA Count Distribution by Column

playerID	0
yearID	0
stint	0
teamID	0
lgID	738
G	0
AB	0
R	0
H	0
2B	0
3B	0
HR	0
RBI	756
SB	2368
CS	23408
BB	0
SO	6341
IBB	36459
HBP	2817
SH	6069
SF	35926
GIDP	25310
birthYear	157
birthMonth	511
birthDay	864
birthCountry	78
birthState	2928
birthCity	412
deathYear	59093
deathMonth	59094
deathDay	59096
deathCountry	59095
deathState	59290
deathCity	59112
nameFirst	38
nameLast	0
nameGiven	38
weight	1143
height	1065
bats	1922
throws	1471
debut	9
finalGame	9
retroID	0
bbrefID	0

```
POS
dtype: int64
```

Dropped 38 observations where nameGiven was not registered

	playerID	yearID	stint	teamID	lgID	G	AB	R	H	2B	\
0	abercda01	1871	1	TRO	NaN	1	4	0	0	0	
1	addybo01	1871	1	RC1	NaN	25	118	30	32	6	
2	addybo01	1873	1	PH2	NaN	10	51	12	16	1	
3	addybo01	1873	2	BS1	NaN	31	152	37	54	6	
4	addybo01	1874	1	HR1	NaN	50	213	25	51	9	

	...	nameFirst	nameLast	nameGiven	weight	\
0	...	Frank	Abercrombie	Francis Patterson	NaN	
1	...	Bob	Addy	Robert Edward	160.0	
2	...	Bob	Addy	Robert Edward	160.0	
3	...	Bob	Addy	Robert Edward	160.0	
4	...	Bob	Addy	Robert Edward	160.0	

	height	bats	throws	debut	POS	FullName
0	NaN	NaN	NaN	1871-10-21	SS	Frank Abercrombie
1	68.0	L	L	1871-05-06	2B	Bob Addy
2	68.0	L	L	1871-05-06	2B	Bob Addy
3	68.0	L	L	1871-05-06	2B	Bob Addy
4	68.0	L	L	1871-05-06	2B	Bob Addy

[5 rows x 34 columns]

```
In [90]: #Merging batting and people on playerID. This gives player metadata
pitching_merged = pd.merge(pitching, people, how='left', on='playerID')
pitching_merged = pd.merge(pitching_merged, positions, on = 'playerID')
pitching_merged = pitching_merged[pitching_merged['POS'] == 'P']

pitching_cols = pitching_merged.columns

#Save csv in working dir
pitching_merged_filename = 'pitching_merged.csv'
pitching_merged.to_csv(pitching_merged_filename)
#Print Logs
print('File saved to {}'.format(current_working_dir))
print('File name: {}'.format(pitching_merged_filename))
print('{} has {} rows and {} columns'.format(pitching_merged_filename,
                                              pitching_merged.shape[0],
                                              pitching_merged.shape[1]))

print('-----')
print('\033[1m' + 'NA Count Distribution by Column' +
      '\033[0m') #Display text output
```



```

print(pitching_merged.isna().sum(axis=0))

#We want to get rid of all of the observations where a given name is not present.
print('-----')
print(
    '\033[1m Dropped {} observations where nameGiven was not registered \033[1m'.
    format(pitching_merged['nameGiven'].isna().sum(axis=0)))
pitching_merged.dropna(subset=['nameGiven'], inplace=True)

#Create a column that merges nameFirst and nameLast
pitching_merged[
    'FullName'] = pitching_merged['nameFirst'] + ' ' + pitching_merged['nameLast']

#Want to drop some unnecessary columns
cols_to_drop = [
    'IBB', 'SH', 'SF', 'GIDP', 'birthMonth', 'birthDay', 'birthState',
    'birthCity', 'deathYear', 'deathMonth', 'deathDay', 'deathCountry',
    'deathCity', 'finalGame', 'retroID', 'bbrefID', 'deathState'
]
print('-----')
print('Dropping Unnecessary Columns')
pitching_merged.drop(columns=cols_to_drop, inplace=True)
print('-----')
print('Displaying Cleaned Pitching Data')
display(pitching_merged.tail(5)) #Display tail of df

```

File saved to C:\Users\jdine\Desktop\SYracuse\Term 7\Scripting for Data Analysis\FinalProject:
File name: pitching_merged.csv
pitching_merged.csv has 43146 rows and 54 columns

----- NA Count Distribution by Column

playerID	0
yearID	0
stint	0
teamID	0
lgID	43
W	0
L	0
G	0
GS	0
CG	0
SHO	0
SV	0
IPouts	0
H	0
ER	0
HR	0
BB	0

```

SO                0
BAOpp             1260
ERA               79
IBB              12364
WP               0
HBP              202
BK               0
BFP              5
GF               0
R                0
SH              16879
SF              16879
GIDP            17999
birthYear        11
birthMonth       48
birthDay         84
birthCountry     7
birthState       835
birthCity        78
deathYear        27299
deathMonth       27299
deathDay         27299
deathCountry     27299
deathState       27344
deathCity        27301
nameFirst        3
nameLast         0
nameGiven        3
weight           299
height           226
bats             480
throws           178
debut            1
finalGame        1
retroID          0
bbrefID          0
POS              0
dtype: int64

```

```

-----
Dropped 3 observations where nameGiven was not registered -----

```

	playerID	yearID	stint	teamID	lgID	W	L	G	GS	CG	\
45801	wilkeaa01	2017	1	MIL	NL	1	0	3	2	0	
45802	willita01	2017	1	MIL	NL	0	0	5	0	0	
45803	woodhu01	2017	1	TBA	AL	0	0	1	0	0	
45804	woodrbr01	2017	1	MIL	NL	2	3	8	8	0	
45805	yacabji01	2017	1	BAL	AL	2	0	14	0	0	

	...	nameFirst	nameLast	nameGiven	weight	height	\
45801	...	Aaron	Wilkerson	Aaron Daniel	190.0	75.0	
45802	...	Taylor	Williams	Taylor Grant	195.0	71.0	
45803	...	Hunter	Wood	Hunter Blake	165.0	73.0	
45804	...	Brandon	Woodruff	Brandon Kyle	215.0	76.0	
45805	...	Jimmy	Yacabonis	James Allin	205.0	75.0	

	bats	throws	debut	POS	FullName
45801	R	R	2017-09-15	P	Aaron Wilkerson
45802	B	R	2017-09-06	P	Taylor Williams
45803	R	R	2017-05-30	P	Hunter Wood
45804	L	R	2017-08-04	P	Brandon Woodruff
45805	R	R	2017-06-11	P	Jimmy Yacabonis

[5 rows x 38 columns]

Writing a program that outputs record holders for desired statistics

- Singe function for batters and pitchers
- Added support for timeslices - Done
- Basic error handling - Done
- Moved column slice out of params/arg
- Maybe add support for by player

```
In [91]: # Want a program that outputs all time records
import numpy as np
```

```
def fetch_records(category, keystatistic, n=5, year=None):
    '''
    category == 'pitching' or 'hitting' str
    key statistic= Filterable stats. Ultimately ranked on this column. str dtype requ
    n = Number of samples to display. Numeric value passed through.
    year = Defaults to none, meaning records are all time. Can set a numeric range be
    '''
    try:
        if 'batting' in category:
            #columns = The columns that run through this program. Defaults to a fixed
            columns = [
                'playerID', 'FullName', 'yearID', 'G', 'AB', 'R', 'H', '2B',
                '3B', 'HR', 'RBI', 'SB', 'CS', 'BB', 'SO'
            ]
            batting_copy = batting_merged.copy() #Create copy of df
            sliced = pd.DataFrame(
                batting_copy[columns]) #Slice df on default cols
            #Introduce Error Handling
            try:
```

```

        if year is not None:
            sliced = sliced[sliced['yearID'] == year]
    except:
        if type(year) != int:
            return (
                'TypeError: Please Enter a Numeric value for the year argument'
            )
        if year > 2017 or year < 1871:
            return (
                'Please enter a valid year between 1871 & 2017, inclusive'
            )

    group = sliced.groupby(['playerID', 'FullName']).sum()
    #Have to manually create a BA column
    group['BattingAverage'] = np.round(group['H'] / group['AB'], 3)
    sort = group.sort_values(by=keystatistic, ascending=False)

    if keystatistic == 'BattingAverage':
        print('Need to reconcile this with an atbat threshold')
    sort.drop(columns=['yearID'], inplace=True)
    sort_indexed = sort[:n]
    #Decoending on parameter entry - These are print logs
    if year is None:
        timeslice = 'Alltime'
    else:
        timeslice = int(year)
    print(
        '\033[1m Displaying top 5 players sorted by {}, Timeslice = {} \033[0m'
        .format(i, timeslice))
    print(
        '-----'
    )
    )
    return sort_indexed
elif 'pitching' in category:
    #columns = The columns that run through this program. Defaults to a fixed
    columns = [
        'playerID', 'yearID', 'W', 'L', 'G', 'GS', 'CG', 'SHO', 'SV',
        'BAOpp', 'ERA', 'FullName'
    ]
    pitching_copy = pitching_merged.copy() #Create copy of df
    pitching_copy.drop_duplicates(inplace=True) #Remove duplicates
    sliced = pd.DataFrame(
        pitching_copy[columns]) #Slice df on default cols
    try:
        if year is not None:
            sliced = sliced[sliced['yearID'] == year]
    except:
        if type(year) != int:

```

```

        return (
            'TypeError: Please Enter a Numeric value for the year argument'
        )
    if year > 2017 or year < 1871:
        return (
            'Please enter a valid year between 1871 & 2017, inclusive'
        )
    #Agg lets us take different measures against different vars in a group
    group = sliced.groupby(['playerID', 'FullName']).agg({
        'W': 'sum',
        'L': 'sum',
        'G': 'sum',
        'GS': 'sum',
        'CG': 'sum',
        'SHO': 'sum',
        'SV': 'sum',
        'BAOpp': 'mean',
        'ERA': 'mean',
    })
    sort = group.sort_values(by=keystatistic, ascending=False)
    sort_indexed = sort[:n]
    if year is None:
        timeslice = 'Alltime'
    else:
        timeslice = int(year)
    print(
        '\033[1m Displaying top 5 players sorted by {}, Timeslice = {} \033[0m'
        .format(i, timeslice))
    print(
        '-----'
    )
    return sort_indexed

except:
    if KeyError:
        print('Please Check Arguments')

```

0.0.4 Testing our Batting Program

```

In [148]: #test output
          keystats = ['H', 'HR', 'RBI']
          #iterate through list of keystats
          for i in keystats:
              display(fetch_records(category='batting',keystatistic=i, n=5, year = 2017))

```

Displaying top 5 players sorted by H, Timeslice = 2017 -----

G	AB	R	H	2B	3B	HR	RBI	SB	CS	\
---	----	---	---	----	----	----	-----	----	----	---

playerID	FullName											
blackch02	Charlie Blackmon	159	644	137	213	35	14	37	104.0	14.0	10.0	
altuvjo01	Jose Altuve	153	590	112	204	39	4	24	81.0	32.0	6.0	
gordode01	Dee Gordon	158	653	114	201	20	9	2	33.0	60.0	16.0	
inciaen01	Ender Inciarte	157	662	93	201	27	5	11	57.0	22.0	9.0	
hosmeer01	Eric Hosmer	162	603	98	192	31	1	25	94.0	6.0	1.0	

		BB	SO	BattingAverage
playerID	FullName			
blackch02	Charlie Blackmon	65	135.0	0.331
altuvjo01	Jose Altuve	58	84.0	0.346
gordode01	Dee Gordon	25	93.0	0.308
inciaen01	Ender Inciarte	49	94.0	0.304
hosmeer01	Eric Hosmer	66	104.0	0.318

Displaying top 5 players sorted by HR, Timeslice = 2017 -----

		G	AB	R	H	2B	3B	HR	RBI	SB	CS	\
playerID	FullName											
stantmi03	Giancarlo Stanton	159	597	123	168	32	0	59	132.0	2.0	2.0	
judgeaa01	Aaron Judge	155	542	128	154	24	3	52	114.0	9.0	4.0	
martijd02	J. D. Martinez	119	432	85	131	26	3	45	104.0	4.0	0.0	
daviskh01	Khrris Davis	153	566	91	140	28	1	43	110.0	4.0	0.0	
gallojo01	Joey Gallo	145	449	85	94	18	3	41	80.0	7.0	2.0	

		BB	SO	BattingAverage
playerID	FullName			
stantmi03	Giancarlo Stanton	85	163.0	0.281
judgeaa01	Aaron Judge	127	208.0	0.284
martijd02	J. D. Martinez	53	128.0	0.303
daviskh01	Khrris Davis	73	195.0	0.247
gallojo01	Joey Gallo	75	196.0	0.209

Displaying top 5 players sorted by RBI, Timeslice = 2017 -----

		G	AB	R	H	2B	3B	HR	RBI	SB	CS	\
playerID	FullName											
stantmi03	Giancarlo Stanton	159	597	123	168	32	0	59	132.0	2.0	2.0	
arenano01	Nolan Arenado	159	606	100	187	43	7	37	130.0	3.0	2.0	
ozunama01	Marcell Ozuna	159	613	93	191	30	2	37	124.0	1.0	3.0	
goldspa01	Paul Goldschmidt	155	558	117	166	34	3	36	120.0	18.0	5.0	
cruzne02	Nelson Cruz	155	556	91	160	28	0	39	119.0	1.0	1.0	

		BB	SO	BattingAverage
playerID	FullName			
stantmi03	Giancarlo Stanton	85	163.0	0.281
arenano01	Nolan Arenado	62	106.0	0.309

ozunama01	Marcell Ozuna	64	144.0	0.312
goldspa01	Paul Goldschmidt	94	147.0	0.297
cruzne02	Nelson Cruz	70	140.0	0.288

0.0.5 Testing our Pitching Program

```
In [95]: keystats = ['W', 'L']
         #iterate through list of keystats
         for i in keystats:
             display(fetch_records(category='pitching',keystatistic=i, n=5, year=None))
```

Displaying top 5 players sorted by W, Timeslice = Alltime -----

playerID	FullName	W	L	G	GS	CG	SHO	SV	BAOpp \
youngcy01	Cy Young	511	316	906	815	749	76	17	0.240000
johnswa01	Walter Johnson	417	279	802	666	531	110	34	0.226667
mathech01	Christy Mathewson	373	188	635	551	434	79	28	0.241176
alexape01	Pete Alexander	373	208	696	599	437	90	32	0.253810
spahnwa01	Warren Spahn	363	245	750	665	382	63	29	0.245909

ERA

playerID	FullName	ERA
youngcy01	Cy Young	2.713043
johnswa01	Walter Johnson	2.348095
mathech01	Christy Mathewson	2.648889
alexape01	Pete Alexander	2.972381
spahnwa01	Warren Spahn	3.277727

Displaying top 5 players sorted by L, Timeslice = Alltime -----

playerID	FullName	W	L	G	GS	CG	SHO	SV	BAOpp	ERA
youngcy01	Cy Young	511	316	906	815	749	76	17	0.240000	2.713043
ryanno01	Nolan Ryan	324	292	807	773	222	61	3	0.208778	3.694074
johnswa01	Walter Johnson	417	279	802	666	531	110	34	0.226667	2.348095
niekrph01	Phil Niekro	318	274	864	716	245	45	29	0.257038	4.098462
perryga01	Gaylord Perry	314	265	777	690	303	53	11	0.253960	3.396000

0.0.6 On to Modeling

- Here we want to merge the hall of fame data with the batting data.
- – Typically train test split. Analyze model/model results
- Show which features are driving hall of fame propensity.

- Build a lightweight function w/ fuzzy matching on user input. User will enter a player name, and the we'll test that player's data using our trained model. Output will be probability dist.

Upsampling minority Class

In [165]: `import sklearn`

```
#This step joins the hall of fame data with the batting data from above. We only need

hof = pd.read_csv(data_links[13]) #read in HOF
hof.head(5)
np.shape(hof)

columns = ['playerID', '2B', '3B', 'AB', 'BAOpp', 'BB', 'CG', 'ERA', \
          'FullName', 'G', 'GS', 'H', 'HR', 'R', 'RBI', 'SB', 'SHO', 'SO', 'SV', 'W'
          ]

merged = pd.concat(
    [batting_merged, pitching_merged], axis=0, ignore_index=True)
merged.drop_duplicates(inplace=True)

merged = pd.DataFrame(merged[columns]) #Only take relevant columns
#Group by for data aggregation
merged = merged.groupby(['playerID', 'FullName']).agg({
    '2B': 'sum',
    '3B': 'sum',
    'AB': 'sum',
    'BAOpp': 'mean',
    'BB': 'sum',
    'CG': 'sum',
    'ERA': 'mean',
    'G': 'sum',
    'GS': 'sum',
    'H': 'sum',
    'HR': 'sum',
    'R': 'sum',
    'RBI': 'sum',
    'SB': 'sum',
    'SHO': 'sum',
    'SO': 'sum',
    'SV': 'sum',
    'W': 'sum',
    'L': 'sum',
})
merged.reset_index(inplace=True)

#Clean/Binarize/Merge HOF data with our position player data
```



```

inducted = hof[hof['inducted'] == 'Y']
players_inducted = inducted[inducted['category'] == 'Player']
players_cols = ['playerID', 'inducted']
players_inducted = players_inducted[players_cols]

hall_merged = pd.merge(merged, players_inducted, how='left', on='playerID')
hall_merged['inducted'].fillna(0, inplace=True)
hall_merged['inducted'].replace('Y', 1, inplace=True)
hall_merged.fillna(0, inplace=True)

#We notice that we have a severe class imbalance problem. Only 230 players are hall of fame.
#stratified sampling to ensure we put an equal dist. of each class in our train test split
#hall_merged[hall_merged['inducted'] == 1].head(5)

#We bring in upsampling techniques to create class balance. In other words, we are randomly sampling from
#our hall of fame class, with replacement, until the number of samples in class equals the number of samples in the
#majority class.
#This prevents us from building a model that randomly guesses the majority class each time.
#https://elitedatascience.com/imbalanced-classes

#We actually shift to downsampling the majority class so that when we predict later, we don't
#just memorize the training data.
df_majority = hall_merged[hall_merged.inducted == 0]
df_minority = hall_merged[hall_merged.inducted == 1]

from sklearn.utils import resample

df_minority_upsampled = resample(
    df_minority,
    replace=True, # sample without replacement
    n_samples=len(df_majority), # to match majority class
    random_state=123) # reproducible results

df_upsampled = pd.concat([df_majority, df_minority_upsampled])

X = df_upsampled.copy()
X.drop(columns=['inducted', 'playerID', 'FullName'], inplace=True)
y = df_upsampled['inducted']

```

Train Test Splits + Grid Searching hyperparams.

```

In [97]: from sklearn.model_selection import train_test_split
from sklearn.grid_search import GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report
import warnings

xtrain, xtest, ytrain, ytest = train_test_split(
    X, y, test_size=0.5, stratify=y)

```

```

#Item freq tells us that we are stratified. We have an equal proportion of each class
from scipy.stats import itemfreq
print('Frequency in Training Set')
print('-----')
print(itemfreq(ytrain))
print('Frequency in Training Set')
print('-----')
print(itemfreq(ytest))

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import GradientBoostingClassifier

model = GradientBoostingClassifier()

#instantiate parm grid
ex_param_grid = {
    "loss": ['deviance', 'exponential'],
    "learning_rate": [0.001, 0.01, 0.1, 1],
    "n_estimators": [100, 200, 300],
    "max_depth": [3, 5, 7, 9]
}

#Grid Search
gsExtC = GridSearchCV(
    model,
    param_grid=ex_param_grid,
    cv=5,
    scoring="accuracy",
    n_jobs=4,
    verbose=1)

#Fit model / Print Score
gsExtC.fit(xtrain, ytrain)
ExtC_best = gsExtC.best_estimator_

```

Frequency in Training Set

```
[[ 0 115]
```

```
[ 1 115]]
```

Frequency in Training Set

```
[[ 0 115]
```

```
[ 1 115]]
```

Fitting 5 folds for each of 96 candidates, totalling 480 fits

```
[Parallel(n_jobs=4)]: Done 52 tasks      | elapsed: 2.8s
```

```
[Parallel(n_jobs=4)]: Done 388 tasks    | elapsed: 9.1s
```

[Parallel(n_jobs=4)]: Done 480 out of 480 | elapsed: 9.8s finished

Showing Our Best Model

```
In [98]: print('-----')
         print('Below is the grid searched best estimator:')
         print('-----')
         print(gsExtC.best_estimator_)
```

```
-----
Below is the grid searched best estimator:
-----
```

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.01, loss='deviance', max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           presort='auto', random_state=None, subsample=1.0, verbose=0,
                           warm_start=False)
```

```
In [99]: print('-----')
         print('Model Score on test set:', np.round(ExtC_best.score(xtest, ytest), 4))
         print('-----')
```

```
-----
Model Score on test set: 0.9043
-----
```

In [276]: *#Here we want to merge our feature importances and features to output feat. ranking*

```
features = list(np.round(clf.feature_importances_, 5))
column_names = ['2B', '3B', 'AB', 'BAOpp', 'BB', 'CG', 'ERA', 'G', 'GS', 'H', 'HR',
                'RBI', 'SB', 'SHO', 'SO', 'SV', 'W', 'L']
```

#merge lists

```
dictionary = dict(zip(column_names, features))
```

```
print('\tReturning Feature Ranking')
```

#Sort Dictionary on values

```
for w in sorted(dictionary, key=dictionary.get, reverse=True):
    print (w, dictionary[w])
```

Returning Feature Ranking

H 0.62934

R 0.28017

```

G 0.03247
RBI 0.01193
HR 0.01179
AB 0.01093
SV 0.01089
BB 0.00344
3B 0.00308
2B 0.00181
SO 0.00152
SB 0.00114
ERA 0.00073
BAOpp 0.00042
GS 0.00036
CG 0.0
SHO 0.0
W 0.0
L 0.0

```

```
In [101]: ypred = ExtC_best.predict(xtest)
```

```

print('Classification Report')
print('-----')
confusion_matrix(ytest, ypred)
print(classification_report(ytest, ypred))

```

```
Classification Report
```

```
-----
```

	precision	recall	f1-score	support
0	0.91	0.90	0.90	115
1	0.90	0.91	0.91	115
avg / total	0.90	0.90	0.90	230

```
In [102]: import itertools
```

```
import matplotlib.pyplot as plt
```

```
#Borrowed from Sklearn Docs
```

```

def plot_confusion_matrix(cm,
                           classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):

```

```
    """
```

```
This function prints and plots the confusion matrix.
```

```

Normalization can be applied by setting `normalize=True`.
"""

plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

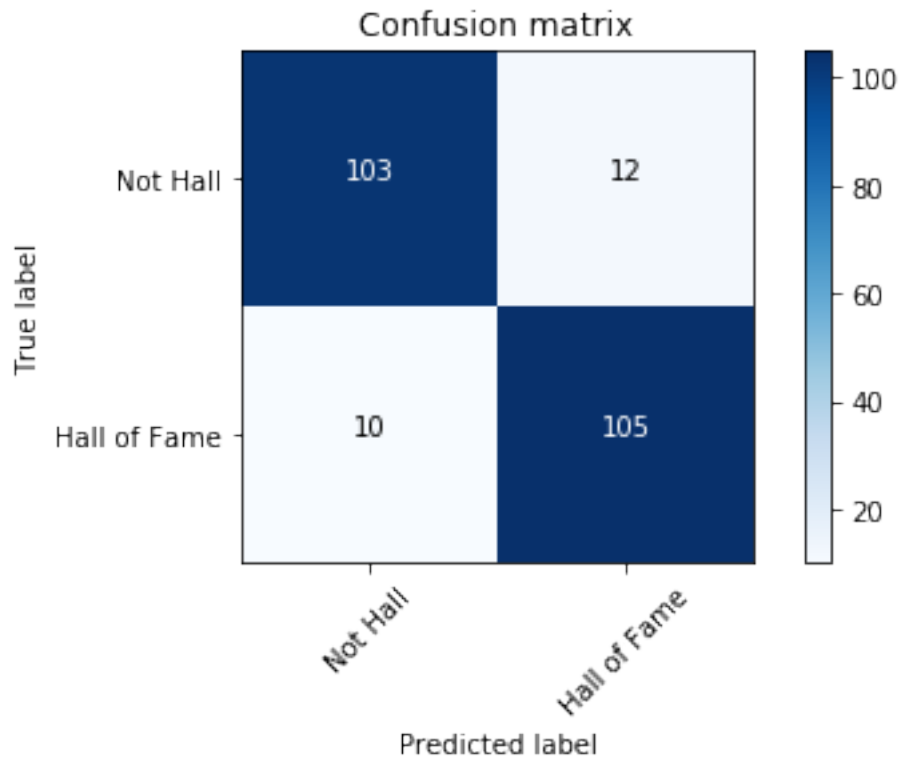
if normalize:
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(
        j,
        i,
        cm[i, j],
        horizontalalignment="center",
        color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')


confusion_mtx = confusion_matrix(ytest, ypred)
# plot the confusion matrix
plot_confusion_matrix(confusion_mtx, classes=['Not Hall', 'Hall of Fame'])

```



Using Pickle to dump Model. This is useful so we don't have retrain our model each time we load this notebook.

```
In [103]: import pickle
           # now you can save it to a file
           with open('ensemble_new.pkl', 'wb') as f:
               pickle.dump(ExtC_best, f)

In [104]: # and later you can load it
           with open('ensemble_new.pkl', 'rb') as f:
               clf = pickle.load(f)
```

Now we want to create a function that takes user input on a player and outputs their probability of making the hall of fame

```
In [149]: #!pip install fuzzywuzzy
           from fuzzywuzzy import fuzz
           from fuzzywuzzy import process
           import time

           #Helper Function
           def fuzzymatching_player(player):
               '''
```

```

input:: User input Movie Title
returns:: Closest Match to user input
'''

name = hall_merged['FullName']
potential_matches = process.extractOne(player, name) #Takes the highest likelihood
return str(potential_matches[0])

def hall_of_fame_projection(userinput=True, player_search=None):
    print(
        '-----'
    )
    print(
        'This program takes a player name as an input, and outputs the players probab
    )
    print(
        '-----'
    )
    try:
        #User input parm
        if userinput == True:
            print('-----')
            player = input('Please Enter a PlayerName:')
            print('-----')
            playername_stringmatch = fuzzymatching_player(player)
            print('-----')
            print('Predicting HOF propensity for {}'.format(
                playername_stringmatch))
            print('-----')
            time.sleep(.1)

        else:
            #Manual parm entry in function instantiation
            player = player_search
            playername_stringmatch = fuzzymatching_player(player)
            print('Predicting HOF propensity for {}'.format(
                playername_stringmatch))
            time.sleep(.1)
        if playername_stringmatch is not None:
            with open('ensemble_new.pkl', 'rb') as f:
                clf = pickle.load(f)
            X = hall_merged.copy()
            X = hall_merged[hall_merged['FullName'] == playername_stringmatch]
            X.drop(columns=['inducted', 'playerID', 'FullName'], inplace=True)
            prob = clf.predict_proba(X)[0][1] * 100
            print('-----')
            print(
                '\n {} has a {}% chance of making the Baseball Hall of Fame,based on

```

```

        format(playername_stringmatch, np.round(prob, 2)))

    print('\t Returning career Statistics through 2017')
    return display(
        hall_merged[hall_merged['FullName'] == playername_stringmatch])

except:
    print('Something went wrong')

```

Using our Model for Inference

Testing chatbox user input

- This works pretty well. Due to the downsampling of the majority class we performed above, this is a more viable representation of how good our model really is, since we aren't training it on the entire dataset.

```
In [150]: hall_of_fame_projection(userinput=True, player_search=None)
```

This program takes a player name as an input, and outputs the players probability of being ele

Please Enter a PlayerName:Jim Abbot

Predicting HOF propensity for Jim Abbott

Jim Abbott has a 80.19% chance of making the Baseball Hall of Fame,based on stats through 201
Returning career Statistics through 2017

	playerID	FullName	2B	3B	AB	BAOpp	BB	CG	ERA	G	\
14	abbotji01	Jim Abbott	0.0	0.0	21.0	0.280364	620	31.0	4.496364	526	
	...	HR	R	RBI	SB	SHO	SO	SV	W	L	inducted
14	...	154	880	3.0	0.0	6.0	898.0	0.0	87.0	108.0	0

[1 rows x 22 columns]

Testing iteration

```
In [151]: players = ['Barry Bonds', 'Cy Young', 'Mike Trout', 'Bryce Harper']
```

```

for i in players:
    hall_of_fame_projection(userinput=False, player_search=i)

```

This program takes a player name as an input, and outputs the players probability of being ele

Predicting HOF propensity for Barry Bonds

Barry Bonds has a 80.62% chance of making the Baseball Hall of Fame,based on stats through 20
Returning career Statistics through 2017

	playerID	FullName	2B	3B	AB	BAOpp	BB	CG	ERA	\		
1552	bondsba01	Barry Bonds	601.0	77.0	9847.0	0.0	2558	0.0	0.0			
	G	...	HR	R	RBI	SB	SHO	SO	SV	W	L	\
1552	2986	...	762	2227	1996.0	514.0	0.0	1539.0	0.0	0.0	0.0	
		inducted										
1552		0										

[1 rows x 22 columns]

This program takes a player name as an input, and outputs the players probability of being ele

Predicting HOF propensity for Cy Young

Cy Young has a 81.53% chance of making the Baseball Hall of Fame,based on stats through 2017
Returning career Statistics through 2017

	playerID	FullName	2B	3B	AB	BAOpp	BB	CG	ERA	\		
18810	youngcy01	Cy Young	87.0	35.0	2960.0	0.24	1298	749.0	2.713043			
	G	...	HR	R	RBI	SB	SHO	SO	SV	W	L	\
18810	1824	...	156	3492	290.0	29.0	76.0	3007.0	17.0	511.0		
		L										
		inducted										
18810	316.0	1										

[1 rows x 22 columns]

This program takes a player name as an input, and outputs the players probability of being ele

Predicting HOF propensity for Mike Trout

Mike Trout has a 46.56% chance of making the Baseball Hall of Fame, based on stats through 2017
 Returning career Statistics through 2017

	playerID	FullName	2B	3B	AB	BAOpp	BB	CG	ERA	G	\
17298	troutmi01	Mike Trout	200.0	40.0	3399.0	0.0	571	0.0	0.0	925	
	...	HR	R	RBI	SB	SHO	SO	SV	W	L	inducted
17298	...	201	692	569.0	165.0	0.0	874.0	0.0	0.0	0.0	0

[1 rows x 22 columns]

 This program takes a player name as an input, and outputs the players probability of being elected

Predicting HOF propensity for Bryce Harper

Bryce Harper has a 35.5% chance of making the Baseball Hall of Fame, based on stats through 2017
 Returning career Statistics through 2017

	playerID	FullName	2B	3B	AB	BAOpp	BB	CG	ERA	G	\
7086	harpebr03	Bryce Harper	149.0	18.0	2756.0	0.0	455	0.0	0.0	768	
	...	HR	R	RBI	SB	SHO	SO	SV	W	L	inducted
7086	...	150	507	421.0	62.0	0.0	665.0	0.0	0.0	0.0	0

[1 rows x 22 columns]

0.0.7 Visualizations

```
In [163]: viz = batting_merged.groupby(['yearID']).sum()
viz.reset_index(inplace=True)
viz
import matplotlib.pyplot as plt
import matplotlib
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (20, 15)
matplotlib.rcParams.update({'font.size': 22})
sns.set_style("darkgrid")
x = np.arange(6)

fig = plt.figure()
```

```

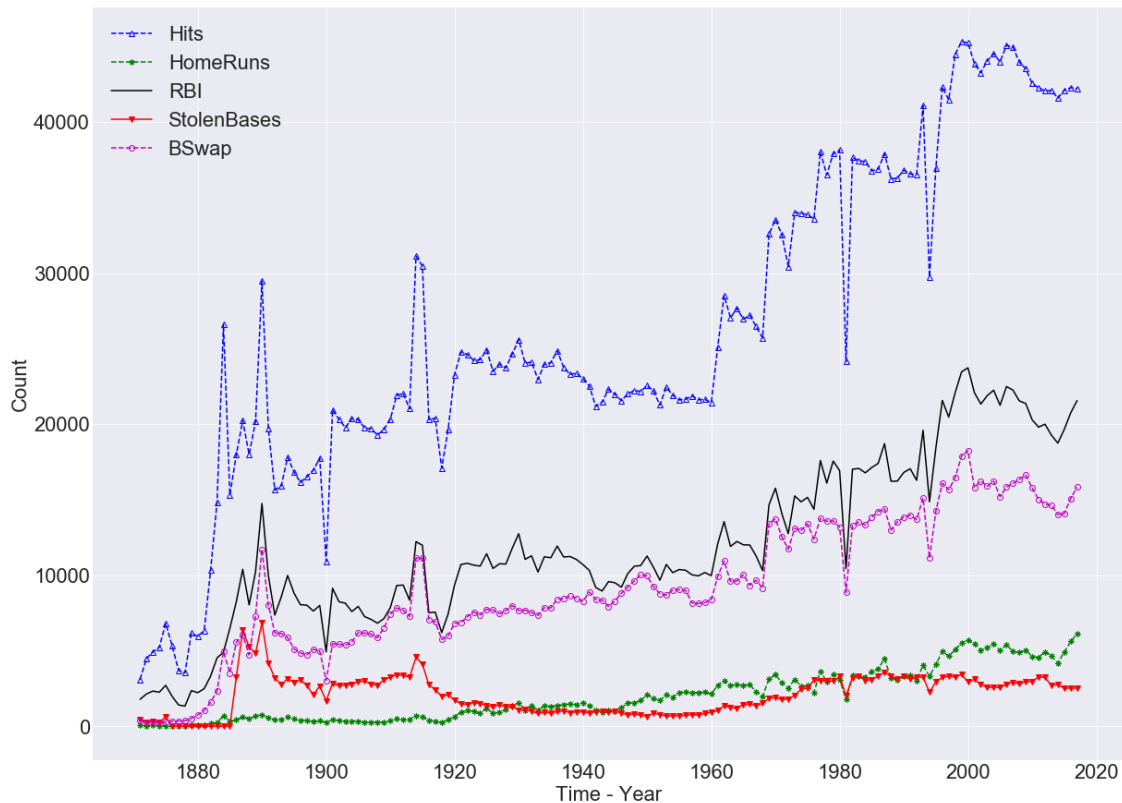
fig.show()
ax = fig.add_subplot(111)

ax.plot(
    viz.yearID,
    viz.H,
    c='b',
    marker="^",
    ls='--',
    label='Hits',
    fillstyle='none')
ax.plot(viz.yearID, viz.HR, c='g', marker=(8, 2, 0), ls='--', label='HomeRuns')
ax.plot(viz.yearID, viz.RBI, c='k', ls='-', label='RBI')
ax.plot(viz.yearID, viz.SB, c='r', marker="v", ls='-', label='StolenBases')
ax.plot(
    viz.yearID,
    viz.BB,
    c='m',
    marker="o",
    ls='--',
    label='BSwap',
    fillstyle='none')

plt.legend(loc=2)
plt.xlabel('Time - Year')
plt.ylabel('Count')
plt.draw()

```

C:\Users\jdine\Documents\Anaconda3\envs\tensorflowlatest\lib\site-packages\matplotlib\figure.py
 "matplotlib is currently using a non-GUI backend, "



```
In [248]: columns = [
            'playerID', 'FullName', 'yearID', 'G', 'AB', 'R', 'H', '2B', '3B', 'HR',
            'RBI', 'SB', 'CS', 'BB', 'SO'
        ]
test = batting_merged.groupby(['playerID', 'FullName']).sum()

test.reset_index(inplace=True)
test = test[columns]

hall_merged = pd.merge(test, players_inducted, how='left', on='playerID')
hall_merged['inducted'].fillna(0, inplace=True)
hall_merged['inducted'].replace('Y', 1, inplace=True)
hall_merged.fillna(0, inplace=True)
hall_merged = hall_merged.groupby(['inducted']).mean()
hall_merged.reset_index(inplace=True)
hall_merged.drop(columns='yearID', inplace=True)

display(hall_merged)
```

	inducted	G	AB	R	H	2B	\
0	0	265.251937	716.501950	92.931713	185.080684	31.375047	
1	1	1673.930435	5854.856522	947.791304	1727.195652	295.178261	

	3B	HR	RBI	SB	CS	BB \
0	6.230510	13.870745	83.447181	14.350735	4.797863	64.727117
1	74.734783	155.691304	874.286957	158.017391	36.847826	643.804348

	SO
0	100.328720
1	583.021739